

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.svm import SVC, LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn import preprocessing
```

```
In [2]: df = pd.read_csv('churn_modelling.csv')
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary      10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [4]: df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], inplace=True)
```

```
In [5]: df.isna().sum()
```

```
Out[5]: CreditScore      0
Geography      0
Gender         0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
```

```
In [6]: X=df.iloc[:, :df.shape[1]-1].values #Independent Variables
y=df.iloc[:, -1].values #Dependent Variable
X.shape, y.shape
```

```
Out[6]: ((10000, 10), (10000,))
```

```
In [7]: print(X[:8,1], '... will now become: ')
label_X_country_encoder = LabelEncoder()
X[:,1] = label_X_country_encoder.fit_transform(X[:,1])
print(X[:8,1])
```

```
['France' 'Spain' 'France' 'France' 'Spain' 'Spain' 'France' 'Germany'] ... will now become:
[0 2 0 0 2 2 0 1]
```

```
In [8]: print(X[:6,2], '... will now become: ')
label_X_gender_encoder = LabelEncoder()
X[:,2] = label_X_gender_encoder.fit_transform(X[:,2])
print(X[:6,2])
```

```
['Female' 'Female' 'Female' 'Female' 'Female' 'Male'] ... will now become:
[0 0 0 0 0 1]
```

```
In [9]: transform = ColumnTransformer([("countries", OneHotEncoder(), [1])], remainder="passthrough")
X = transform.fit_transform(X)
X
```

```
Out[9]: array([[1.0, 0.0, 0.0, ..., 1, 1, 101348.88],
               [0.0, 0.0, 1.0, ..., 0, 1, 112542.58],
               [1.0, 0.0, 0.0, ..., 1, 0, 113931.57],
               ...,
               [1.0, 0.0, 0.0, ..., 0, 1, 42085.58],
               [0.0, 1.0, 0.0, ..., 1, 0, 92888.52],
               [1.0, 0.0, 0.0, ..., 1, 0, 38190.78]], dtype=object)
```

```
In [10]: X = X[:,1:]
X.shape
```

```
Out[10]: (10000, 11)
```

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
In [12]: sc=StandardScaler()
X_train[:,np.array([2,4,5,6,7,10])] = sc.fit_transform(X_train[:,np.array([2,4,5,6,7,10])])
X_test[:,np.array([2,4,5,6,7,10])] = sc.transform(X_test[:,np.array([2,4,5,6,7,10])])
```

```
In [14]: sc=StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
X_train
```

```
Out[14]: array([[ -0.5698444 ,  1.74309049,  0.16958176, ...,  0.64259497,
        -1.03227043,  1.10643166],
       [ 1.75486502, -0.57369368, -2.30455945, ...,  0.64259497,
        0.9687384 , -0.74866447],
       [-0.5698444 , -0.57369368, -1.19119591, ...,  0.64259497,
        -1.03227043,  1.48533467],
       ...,
       [-0.5698444 , -0.57369368,  0.9015152 , ...,  0.64259497,
        -1.03227043,  1.41231994],
       [-0.5698444 ,  1.74309049, -0.62420521, ...,  0.64259497,
        0.9687384 ,  0.84432121],
       [ 1.75486502, -0.57369368, -0.28401079, ...,  0.64259497,
        -1.03227043,  0.32472465]])
```

```
In [15]: from tensorflow.keras.models import Sequential
# Initializing the ANN
classifier = Sequential()
```

```
In [16]: from tensorflow.keras.layers import Dense
```

```
In [18]: classifier.add(Dense(activation = 'relu', input_dim = 11, units=256, kernel_initializer='unifo
```

```
In [19]: classifier.add(Dense(activation = 'relu', units=512, kernel_initializer='uniform'))
classifier.add(Dense(activation = 'relu', units=256, kernel_initializer='uniform'))
classifier.add(Dense(activation = 'relu', units=128, kernel_initializer='uniform'))
```

```
In [20]: classifier.add(Dense(activation = 'sigmoid', units=1, kernel_initializer='uniform'))
```

```
In [21]: classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [22]: classifier.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	3072
dense_2 (Dense)	(None, 512)	131584
dense_3 (Dense)	(None, 256)	131328
dense_4 (Dense)	(None, 128)	32896
dense_5 (Dense)	(None, 1)	129

=====
Total params: 299009 (1.14 MB)
Trainable params: 299009 (1.14 MB)
Non-trainable params: 0 (0.00 Byte)

```
In [23]: classifier.fit(  
    X_train, y_train,  
    validation_data=(X_test,y_test),  
    epochs=20,  
    batch_size=32  
)
```

```
Epoch 1/20
250/250 [=====] - 6s 12ms/step - loss: 0.4291 - accuracy: 0.8200 - v
al_loss: 0.3551 - val_accuracy: 0.8615
Epoch 2/20
250/250 [=====] - 2s 10ms/step - loss: 0.3606 - accuracy: 0.8553 - v
al_loss: 0.3629 - val_accuracy: 0.8700
Epoch 3/20
250/250 [=====] - 3s 11ms/step - loss: 0.3450 - accuracy: 0.8584 - v
al_loss: 0.3508 - val_accuracy: 0.8575
Epoch 4/20
250/250 [=====] - 3s 11ms/step - loss: 0.3416 - accuracy: 0.8629 - v
al_loss: 0.3407 - val_accuracy: 0.8625
Epoch 5/20
250/250 [=====] - 3s 11ms/step - loss: 0.3387 - accuracy: 0.8611 - v
al_loss: 0.3357 - val_accuracy: 0.8675
Epoch 6/20
250/250 [=====] - 3s 11ms/step - loss: 0.3356 - accuracy: 0.8656 - v
al_loss: 0.3392 - val_accuracy: 0.8650
Epoch 7/20
250/250 [=====] - 3s 10ms/step - loss: 0.3309 - accuracy: 0.8649 - v
al_loss: 0.3372 - val_accuracy: 0.8580
Epoch 8/20
250/250 [=====] - 3s 10ms/step - loss: 0.3292 - accuracy: 0.8664 - v
al_loss: 0.3379 - val_accuracy: 0.8625
Epoch 9/20
250/250 [=====] - 3s 11ms/step - loss: 0.3262 - accuracy: 0.8670 - v
al_loss: 0.3729 - val_accuracy: 0.8485
Epoch 10/20
250/250 [=====] - 3s 11ms/step - loss: 0.3236 - accuracy: 0.8706 - v
al_loss: 0.3482 - val_accuracy: 0.8515
Epoch 11/20
250/250 [=====] - 3s 11ms/step - loss: 0.3182 - accuracy: 0.8710 - v
al_loss: 0.3340 - val_accuracy: 0.8675
Epoch 12/20
250/250 [=====] - 3s 11ms/step - loss: 0.3155 - accuracy: 0.8696 - v
al_loss: 0.3402 - val_accuracy: 0.8595
Epoch 13/20
250/250 [=====] - 3s 10ms/step - loss: 0.3102 - accuracy: 0.8742 - v
al_loss: 0.3418 - val_accuracy: 0.8620
Epoch 14/20
250/250 [=====] - 3s 10ms/step - loss: 0.3068 - accuracy: 0.8730 - v
al_loss: 0.3674 - val_accuracy: 0.8580
Epoch 15/20
250/250 [=====] - 3s 11ms/step - loss: 0.3036 - accuracy: 0.8756 - v
al_loss: 0.3448 - val_accuracy: 0.8580
Epoch 16/20
250/250 [=====] - 3s 11ms/step - loss: 0.2988 - accuracy: 0.8795 - v
al_loss: 0.3455 - val_accuracy: 0.8595
Epoch 17/20
250/250 [=====] - 3s 11ms/step - loss: 0.2915 - accuracy: 0.8838 - v
al_loss: 0.3696 - val_accuracy: 0.8515
Epoch 18/20
250/250 [=====] - 3s 11ms/step - loss: 0.2837 - accuracy: 0.8855 - v
al_loss: 0.3598 - val_accuracy: 0.8590
Epoch 19/20
250/250 [=====] - 3s 10ms/step - loss: 0.2813 - accuracy: 0.8845 - v
al_loss: 0.3747 - val_accuracy: 0.8565
Epoch 20/20
250/250 [=====] - 3s 10ms/step - loss: 0.2740 - accuracy: 0.8882 - v
al_loss: 0.3831 - val_accuracy: 0.8600
```

Out[23]: <keras.src.callbacks.History at 0x19347b54150>

```
In [32]: y_pred = classifier.predict(X_test)
y_pred
```

63/63 [=====] - 0s 5ms/step

```
Out[32]: array([[0.40783623],
               [0.2887754 ],
               [0.10280854],
               ...,
               [0.09997811],
               [0.19579199],
               [0.2798526 ]], dtype=float32)
```

```
In [33]: y_pred = (y_pred > 0.5)
y_pred
```

```
Out[33]: array([[False],
               [False],
               [False],
               ...,
               [False],
               [False],
               [False]])
```

```
In [34]: from sklearn.metrics import confusion_matrix, classification_report
```

```
In [38]: cm1 = confusion_matrix(y_test, y_pred)
```

```
In [39]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.91	1595
1	0.69	0.56	0.62	405
accuracy			0.86	2000
macro avg	0.79	0.75	0.77	2000
weighted avg	0.85	0.86	0.85	2000

```
In [40]: accuracy_model1 = ((cm1[0][0]+cm1[1][1])*100)/(cm1[0][0]+cm1[1][1]+cm1[0][1]+cm1[1][0])
print (accuracy_model1, '% of testing data was classified correctly')
```

86.0 % of testing data was classified correctly

```
In [ ]:
```